Predicting gains beyond shared token selection

Held-out images and class pairs · frozen CLIP ViT-B/16

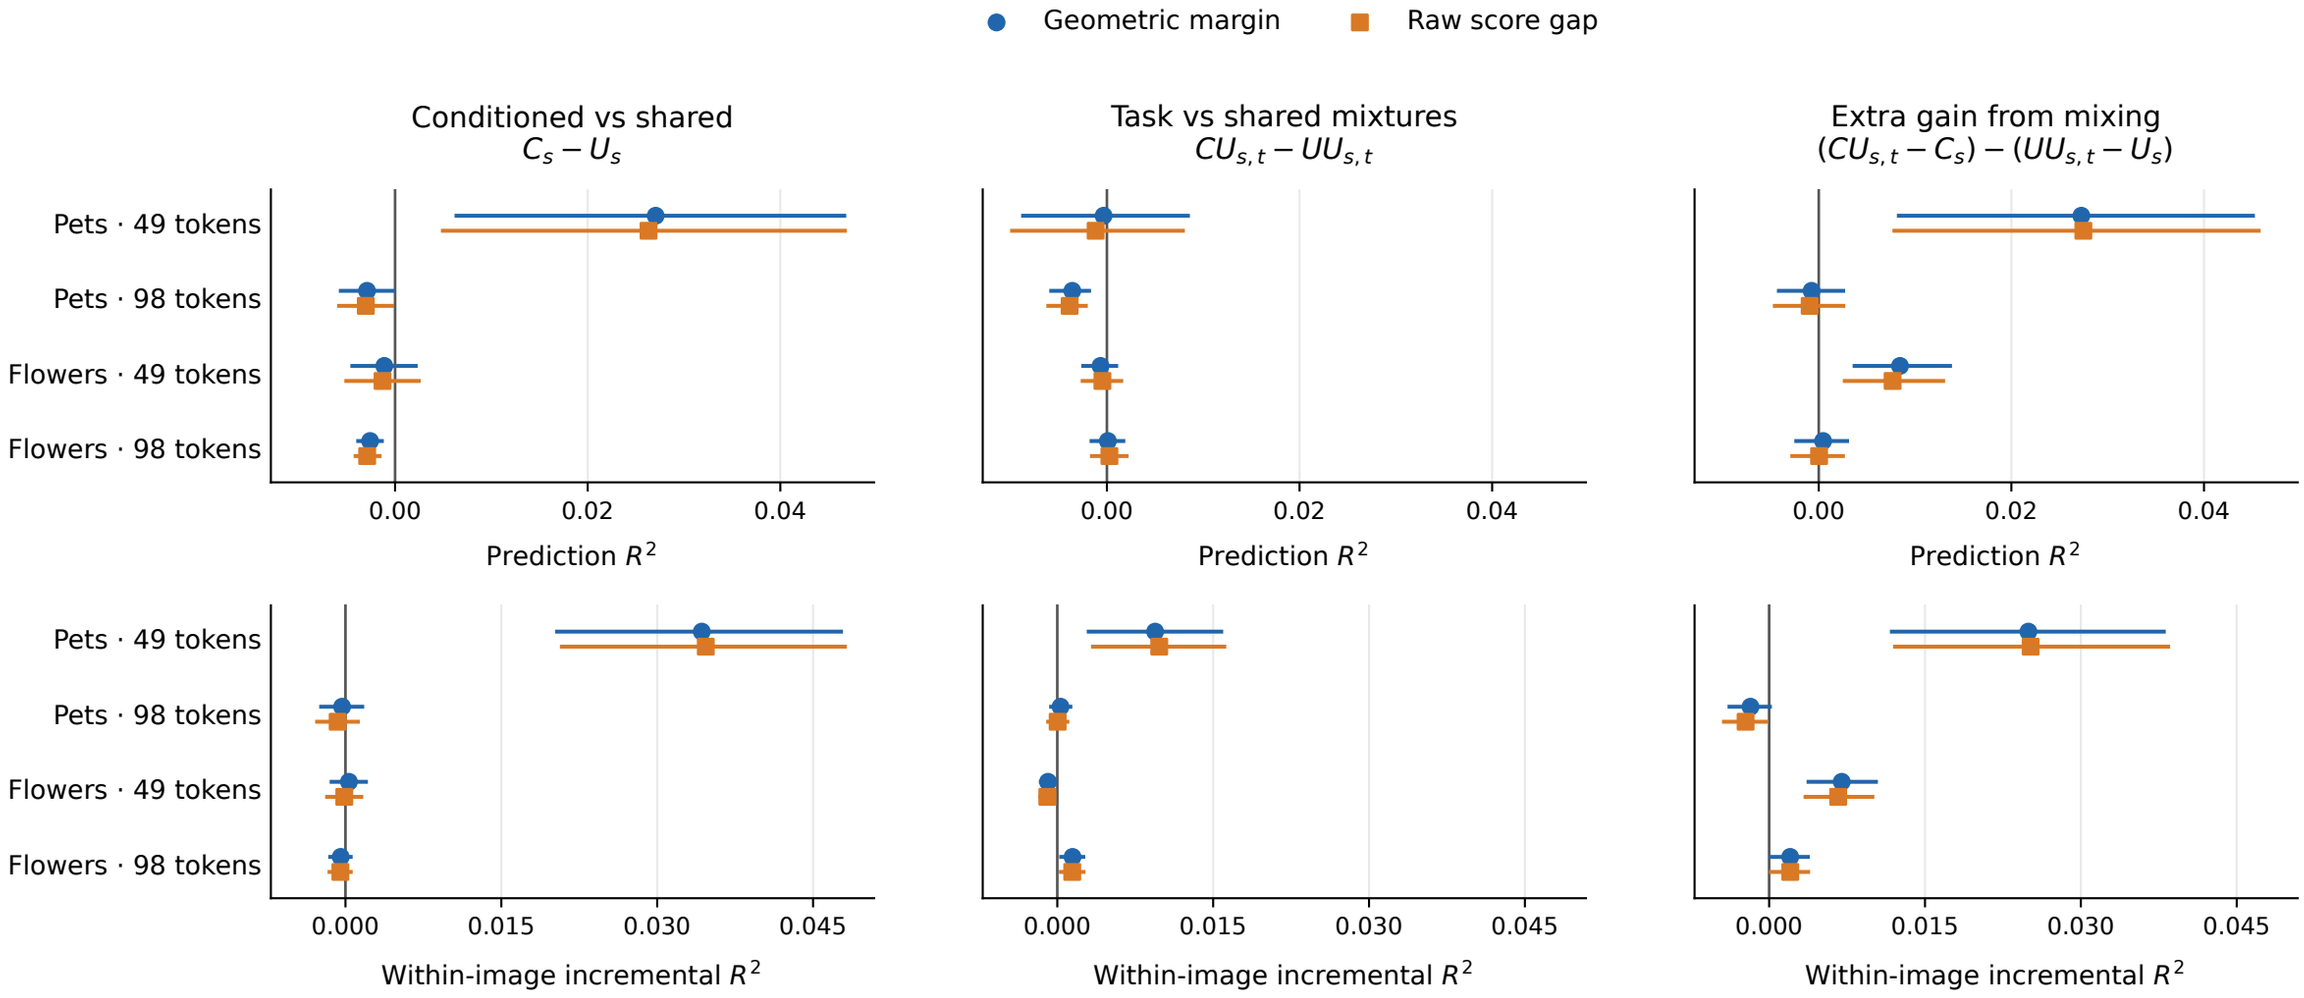

Points pool three seeds (left) or six ordered seed pairs (middle/right); bars are ordinary 95% image-bootstrap CIs.
Top: prediction error relative to the training-mean predictor. Negative values indicate worse prediction.
Bottom: added value of selection features beyond text distance and shared boundary gap, after centering within each image.
A positive increment does not imply positive absolute prediction R².