How pruning depth changes the value of conditioning

Frozen CLIP ViT-B/16 · 49 retained patch tokens · fixed training protocol

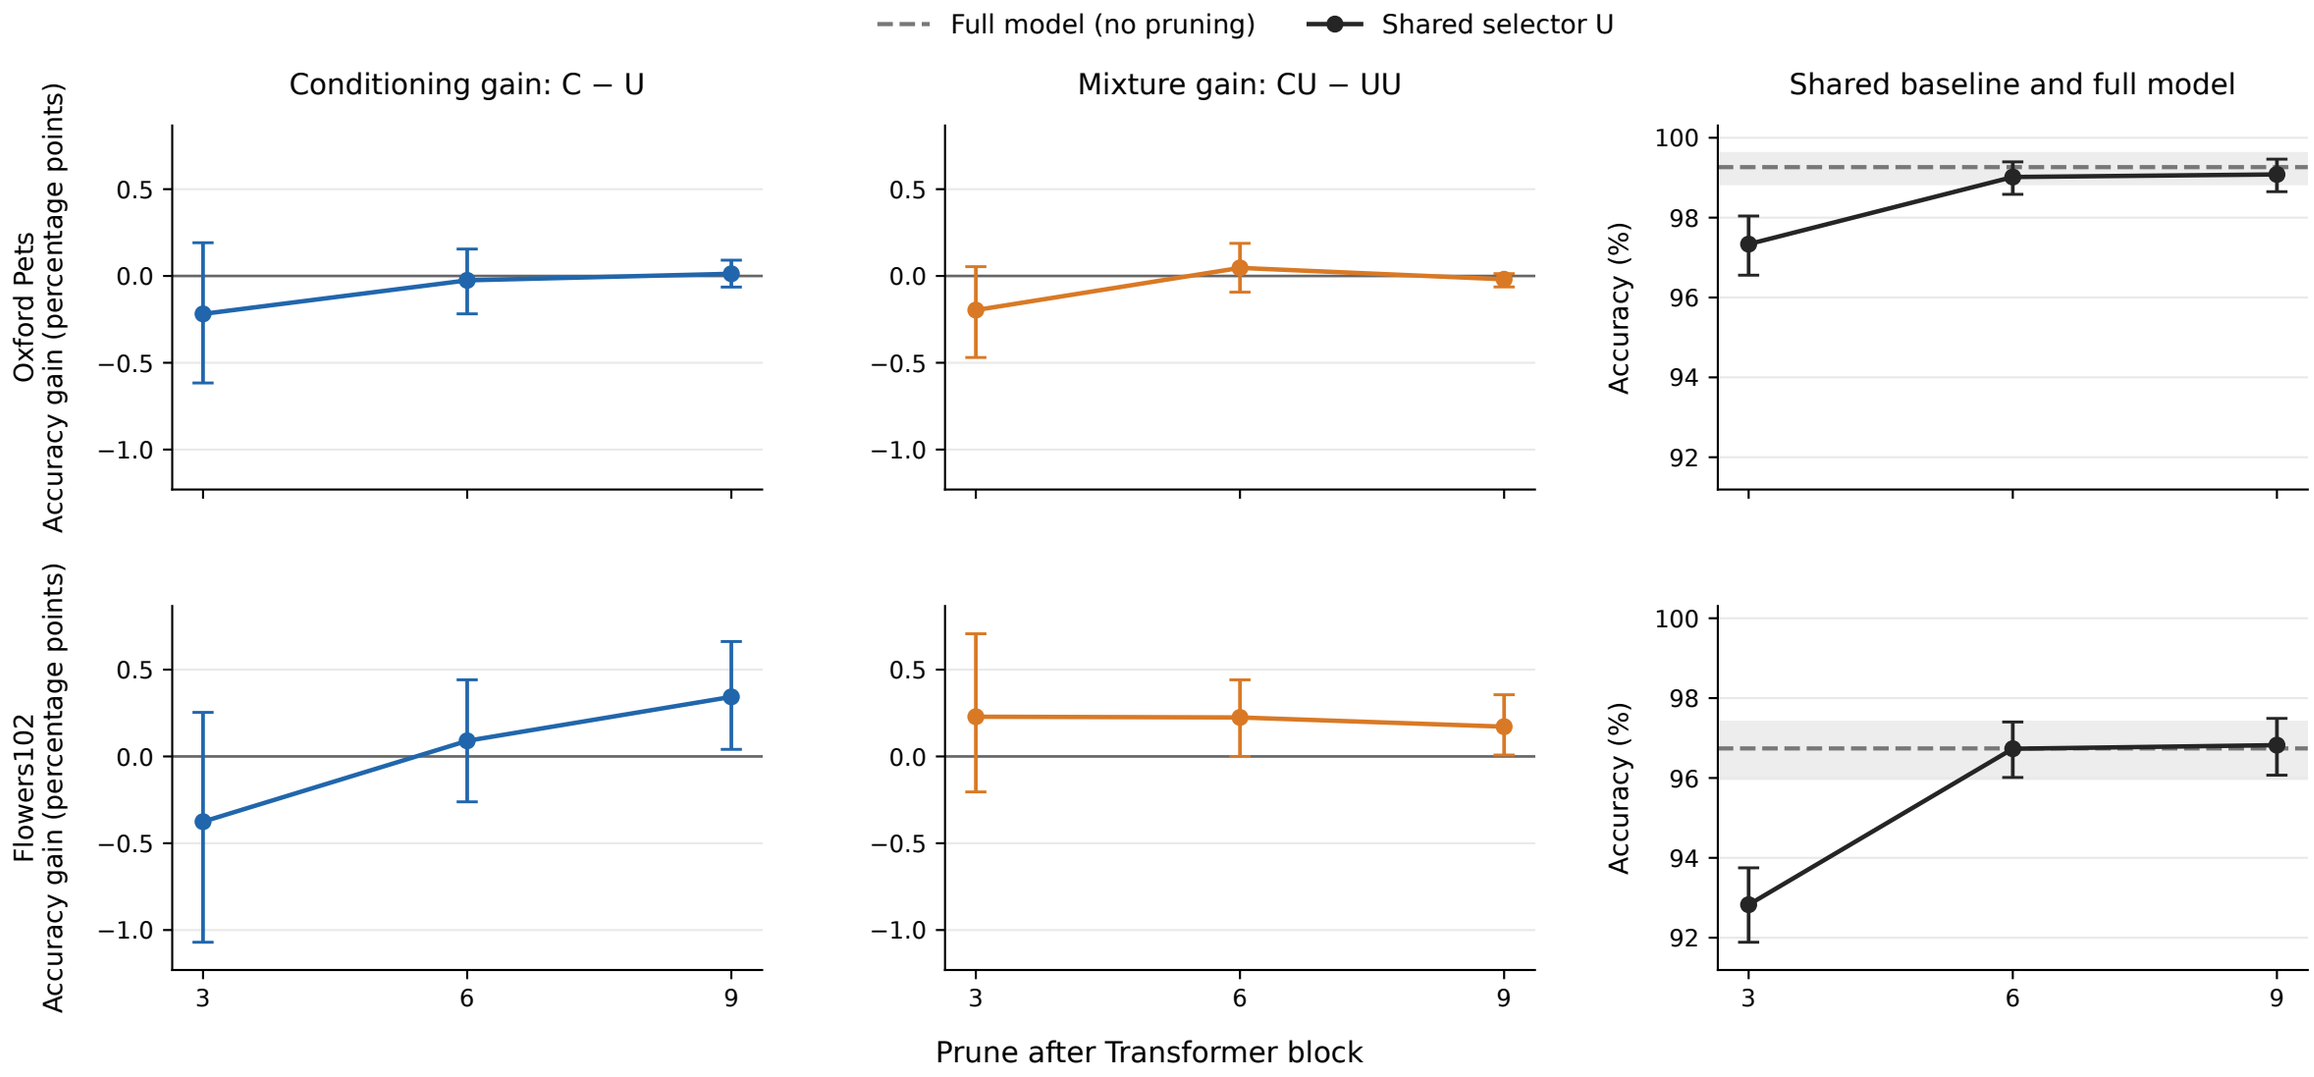

C − U pools three paired seeds; CU − UU pools six ordered seed pairs. Bars show ordinary 95% image-bootstrap intervals.
Each depth uses its own fixed-step selectors and teacher; images, tasks, teacher probe masks and training settings are held fixed.
The right column shows the shared selector's absolute performance against the unchanged full model.
Depth effects are quantified separately by paired changes relative to block 6, with family-adjusted intervals.