Same image, different candidate classes

(a) Controlled inputs

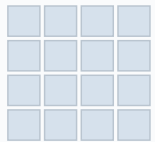

one image x

Frozen ViT prefix

patch + CLS features H

Change the competitor

$T_A = \{a, b\}$ $T_B = \{a, c\}$

frozen text → descriptors d_A, d_B

(b) Paired token selection

same H

d_A

$C(H, d_A)$ → S_A

same trained selector

d_B

$C(H, d_B)$ → S_B

same k

$U(H)$ → S_U

shared across tasks

C: 348k parameters U: 347k parameters

(c) Three value tests

1 Use Exchange selections

	S_A	S_B
T_A	own	swapped
T_B	swapped	own

2 Cost Reallocate selector capacity

C vs. wider U

comparable parameter budgets

3 Value Mixing, budget, depth

$C_s + U_t$ vs. $U_s + U_t$

same donor and retained-token budget

Evaluate retained tokens through the same frozen ViT suffix and CLIP classifier; test unseen class combinations.