Does directly learning the retained set improve classification?

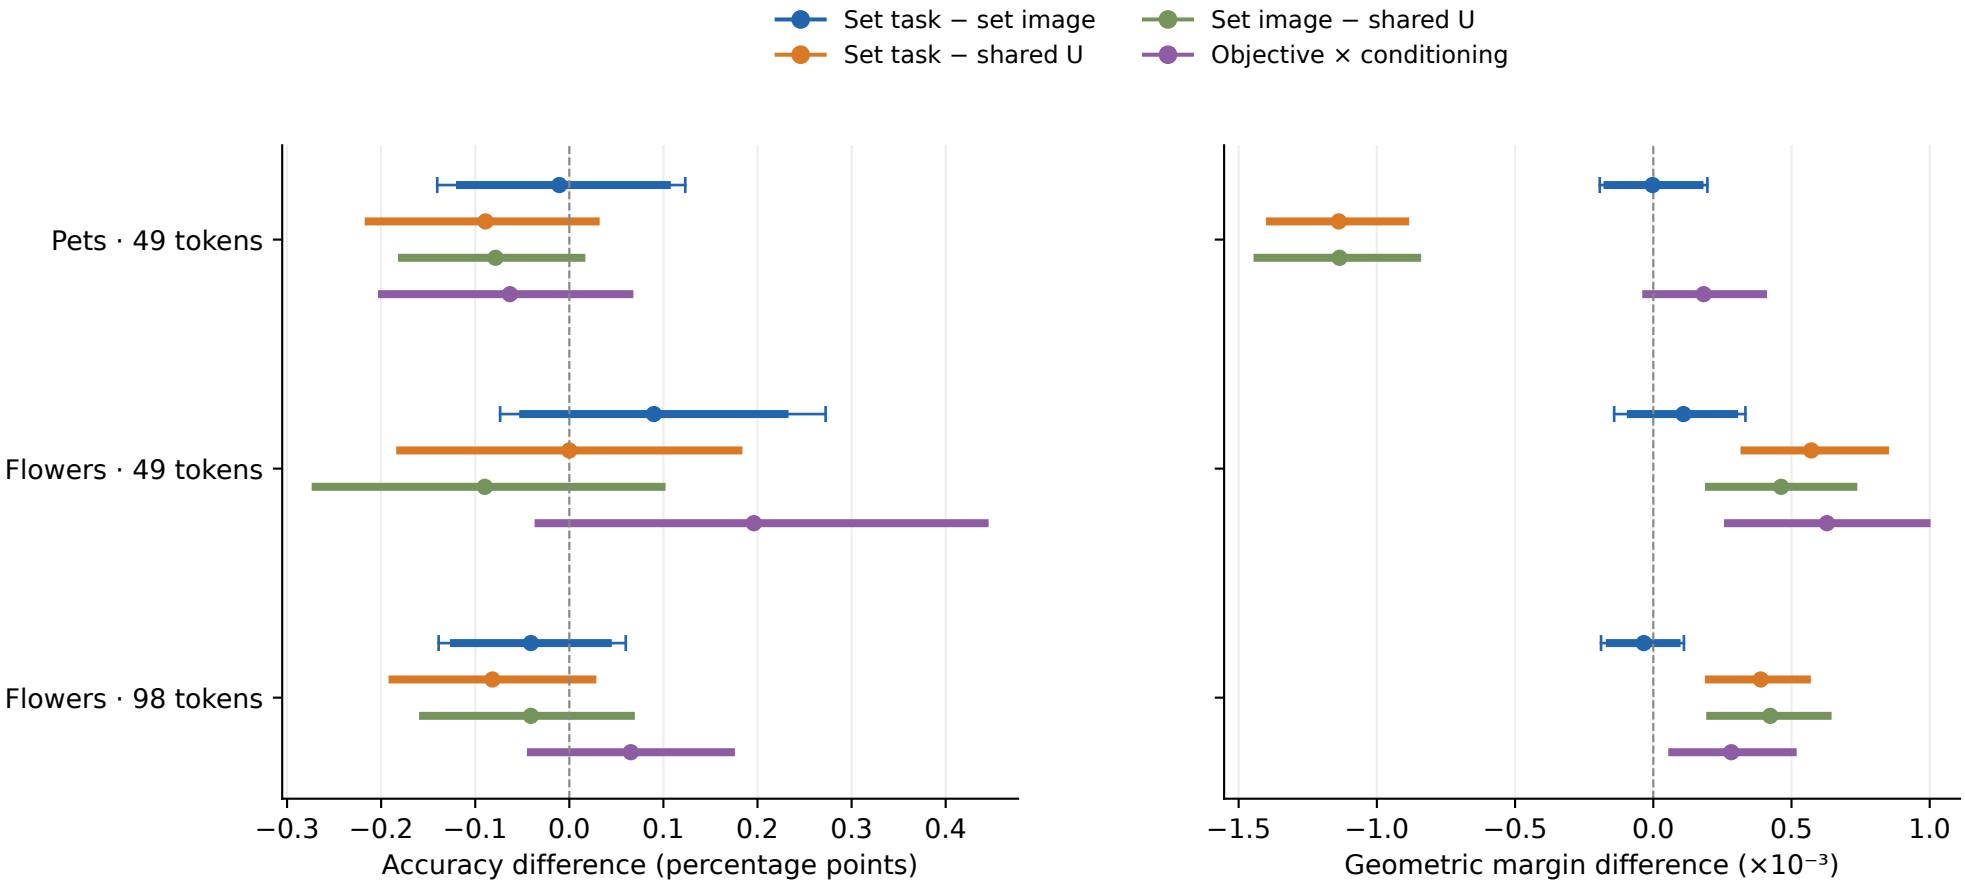

Same fixed hybrid masks; matched task/image branches; all three seeds retained.
Thick bars: paired image-bootstrap 95% intervals. Thin blue: three-setting adjustment.
Interaction = (task − image) with direct-set loss minus (task − image) with KL.